

# Floyd-Warshall Algorithm for All Pair Shortest Path

A dark blue, abstract, curved shape that starts from the bottom left and extends diagonally upwards towards the right, filling the bottom half of the slide.

# OVERVIEW:

- Problem Statement
- Sequential Algorithm
- Rectangular Algorithm
- Blocked Floyd Warshall
- Implementation and Parallel Algorithm
- Cache Optimized Algorithm
- Results
- References

# The Floyd Warshall Algorithm

- This algorithm solves the shortest path problem for a directed and weighted graph.
- It tries to find the minimum distance between any pair of vertices in a graph. In the adjacency matrix created  $\text{dist}(i,j)$  represents the distance from source (ith node) to the destination(jth node).
- Since the distance from a vertex to itself is going to be 0 , hence all the diagonals are set to 0 in the matrix.
- In this we consider every vertex as an intermediate vertex 'k' and find if the distance between i,j through k is smaller than the existing distance.  
i.e.  $\text{dist}(i,j) = \min( \text{dist}(i,j) , \text{dist}(i,k) + \text{dist}(k,j) )$  .

# SEQUENTIAL ALGORITHM

Input:  $n$  = number of vertices

$A$  = adjacency matrix

Output: Transformed  $A$  that contains all pair shortest paths

For each vertex  $k$  from 1 to  $n$ :

For each pair of vertices  $i$  and  $j$  from 1 to  $n$ :

If ( $\text{distance}[i][k] \neq \text{INF} \ \&\& \ \text{distance}[k][j] \neq \text{INF}$ ) {

    If  $\text{dist}[i][j] > \text{dist}[i][k] + \text{dist}[k][j]$ , then update  $\text{dist}[i][j]$  to  $\text{dist}[i][k] + \text{dist}[k][j]$

end for

end for

- Clearly it can be observed that the time complexity of this algorithm is  $O(n^3)$ .

# Rectangular Implementation –

- The Rectangular algorithm efficiently solves the same problem as Floyd-Warshall by skipping iterations where the distance between vertices is infinite, reducing unnecessary computations for faster execution.
- Clearly the time complexity is  $O(n^3)$ .

Input: num\_vertices, dist (adjacency matrix of size num\_vertices x num\_vertices)

Output: Updated adjacency matrix containing all-pair shortest paths

```
procedure solve(dist, num_vertices)
  for j from 0 to num_vertices - 1
    for i from 0 to num_vertices - 1
      r ← dist[i][j]
      if r equals INF
        continue
      for k from 0 to num_vertices - 1
        t ← dist[j][k]
        if t equals INF or i equals j or k equals j
          or i equals k
          continue
        s ← dist[i][k]
        if s is greater than r + t
          s ← r + t
```

# PARALLEL BLOCKED ALGORITHM

- Divide the distance matrix into smaller blocked matrices to improve cache locality, accelerating memory operations.
- The original  $n \times n$  matrix  $A$  is divided into smaller  $b \times b$  sub-matrices ( $A_{ij}$ ,  $A_{ik}$ ,  $A_{kj}$ ,  $A_{kk}$ ), where each sub-matrix has an optimal blocking factor  $b$ . Determining this factor is crucial for cache efficiency, as it ensures that the adjacency matrix fits within cache constraints.
- Our computation has been performed in the OpenMP environment.

This blocked parallel implementation, proposed must be processed in three separate phases:

- Dependent phase – processing the  $k$ th diagonal block.
- Partially dependent phase – processing the  $k$ th row and the  $k$ th column of blocks.
- Independent phase – processing the remaining blocks.

Image source: [3] Towards performance improvement of a parallel Floyd-Warshall algorithm using OpenMP and Intel TBB

|           |           |           |           |           |           |
|-----------|-----------|-----------|-----------|-----------|-----------|
| $W_{1,1}$ | $W_{1,2}$ | $W_{1,3}$ | $W_{1,4}$ | $W_{1,5}$ | $W_{1,6}$ |
| $W_{2,1}$ | $W_{2,2}$ | $W_{2,3}$ | $W_{2,4}$ | $W_{2,5}$ | $W_{2,6}$ |
| $W_{3,1}$ | $W_{3,2}$ | $W_{3,3}$ | $W_{3,4}$ | $W_{3,5}$ | $W_{3,6}$ |
| $W_{4,1}$ | $W_{4,2}$ | $W_{4,3}$ | $W_{4,4}$ | $W_{4,5}$ | $W_{4,6}$ |
| $W_{5,1}$ | $W_{5,2}$ | $W_{5,3}$ | $W_{5,4}$ | $W_{5,5}$ | $W_{5,6}$ |
| $W_{6,1}$ | $W_{6,2}$ | $W_{6,3}$ | $W_{6,4}$ | $W_{6,5}$ | $W_{6,6}$ |

(a) Dependent phase

|           |           |           |           |           |           |
|-----------|-----------|-----------|-----------|-----------|-----------|
| $W_{1,1}$ | $W_{1,2}$ | $W_{1,3}$ | $W_{1,4}$ | $W_{1,5}$ | $W_{1,6}$ |
| $W_{2,1}$ | $W_{2,2}$ | $W_{2,3}$ | $W_{2,4}$ | $W_{2,5}$ | $W_{2,6}$ |
| $W_{3,1}$ | $W_{3,2}$ | $W_{3,3}$ | $W_{3,4}$ | $W_{3,5}$ | $W_{3,6}$ |
| $W_{4,1}$ | $W_{4,2}$ | $W_{4,3}$ | $W_{4,4}$ | $W_{4,5}$ | $W_{4,6}$ |
| $W_{5,1}$ | $W_{5,2}$ | $W_{5,3}$ | $W_{5,4}$ | $W_{5,5}$ | $W_{5,6}$ |
| $W_{6,1}$ | $W_{6,2}$ | $W_{6,3}$ | $W_{6,4}$ | $W_{6,5}$ | $W_{6,6}$ |

(b) Partially dependent phase

|           |           |           |           |           |           |
|-----------|-----------|-----------|-----------|-----------|-----------|
| $W_{1,1}$ | $W_{1,2}$ | $W_{1,3}$ | $W_{1,4}$ | $W_{1,5}$ | $W_{1,6}$ |
| $W_{2,1}$ | $W_{2,2}$ | $W_{2,3}$ | $W_{2,4}$ | $W_{2,5}$ | $W_{2,6}$ |
| $W_{3,1}$ | $W_{3,2}$ | $W_{3,3}$ | $W_{3,4}$ | $W_{3,5}$ | $W_{3,6}$ |
| $W_{4,1}$ | $W_{4,2}$ | $W_{4,3}$ | $W_{4,4}$ | $W_{4,5}$ | $W_{4,6}$ |
| $W_{5,1}$ | $W_{5,2}$ | $W_{5,3}$ | $W_{5,4}$ | $W_{5,5}$ | $W_{5,6}$ |
| $W_{6,1}$ | $W_{6,2}$ | $W_{6,3}$ | $W_{6,4}$ | $W_{6,5}$ | $W_{6,6}$ |

(c) Independent phase

# ARCHITECTURE

The following implementations of the FW algorithm were executed in C++ using OpenMP environments.

|                       |                       |
|-----------------------|-----------------------|
| Processor Name        | Intel Xeon E5-2640 v4 |
| Number of Cores       | 20                    |
| Number of Threads     | 40                    |
| Base Core Clocks[GHz] | 2.40                  |
| L1 Cache[KB]          | 32                    |
| L1 Cache Line(Byte)   | 64                    |



# Optimal Blocking Factor

According to Pozder, Nudžejma & Ćorović, Dalila & Herenda, Esma & Divjan, Belmin. (2021)[3],

The optimal blocking factor  $b$  must be the largest integer such that  $2 \cdot b^2 \cdot 4 \leq C$  and  $b$  is a dividable of  $S/4$ , where  $C$  presents L1 cache capacity and  $S$  is L1 cache line.

Our system has  $C = 32$  KB and  $S = 64$  bytes

On solving, we obtain that  $b \leq 63.24$  and  $b$  should be divisible by 16.

Taking these two factors into consideration we consider  $b = 16$

# THE ALGORITHM

```
1: floyd(C,A,B) {
2: C, A and B are b×b matrices
3: for k from 0 to b do
4:   for j from 0 to b do
5:     for i from 0 to b do
6:        $C[i][j] = \min(C[i][j], A[i][k] + B[k][j])$ 
7:     end for
8:   end for
9: end for
10: }
11: blocked_floyd_warshall(W,n) {
12: //split W into "blocks" with blocksize "b"
13: //For simplicity if b divides n, b blocks are created
14: for k from 0 to B do
15:   //Dependent phase
16:   floyd(Wkk, Wkk, Wkk);
17:   //Partially Dependent Phase
18:   parallel for j from 0 and j ≠ k to B
19:     floyd(Wkj , Wkk, Wkj );
20:   end parallel for
21:   parallel for i from 0 and i ≠ k to B
22:     floyd(Wik, Wik, Wkk);
23:   for j from 0 and j ≠ k to B do
24:     //Independent Phase
25:     floyd(Wij , Wik, Wkj );
26:   end for
27: end parallel for
28: end for
29: }
```

REFERENCE: [3]

## Towards performance improvement of a parallel Floyd-Warshall algorithm using OpenMP and Intel TBB

- Here the 2D distance matrix is converted into a 1D array because in C++ memory is allocated in contiguous blocks in row major fashion also contributing to the creation of a cache efficient algorithm as elements are accessed in a nested loop fashion and parallelization adds on to the overall reduction in time.

# CACHE OPTIMIZED ALGORITHM

- This algorithm is similar to the blocked implementation but it has been done without the parallelization to observe how efficient the Floyd Warshall algorithm can be if it is cache optimized.
- In C++, converting a 2D distance matrix into a 1D array is done to leverage memory allocation in contiguous blocks, following the row major fashion. This enhances cache efficiency, especially when elements are accessed in nested loop fashion.

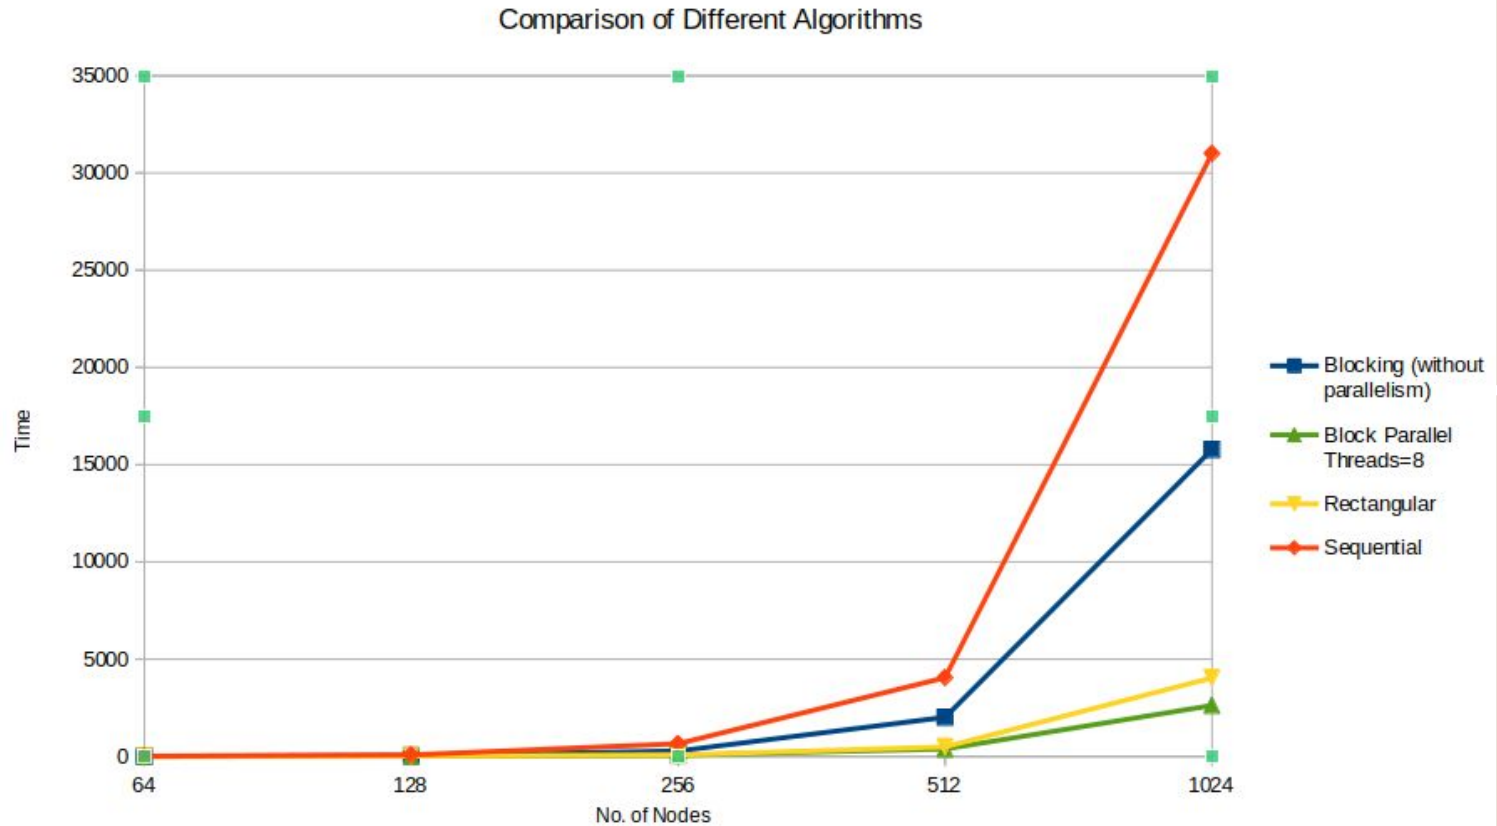
The codes for all the algorithms can be viewed here:

<https://github.com/b22cs005/Optimization-of-Floyd-Warshall-Algorithm>

# RESULTS

COMPARING ALL THE IMPLEMENTED ALGORITHMS:

| Nodes | Sequential | Rectangular | Cache Optimized | Blocked Parallel |
|-------|------------|-------------|-----------------|------------------|
| 32    | 1          | 0           | 1               | 1                |
| 64    | 10         | 1           | 6               | 3                |
| 128   | 79         | 10          | 45              | 9                |
| 256   | 654        | 77          | 277             | 56               |
| 512   | 4049       | 491         | 2015            | 381              |
| 1024  | 31000      | 4040        | 15780           | 2620             |
| 2048  | 248196     | 33247       | 127250          | 19548            |

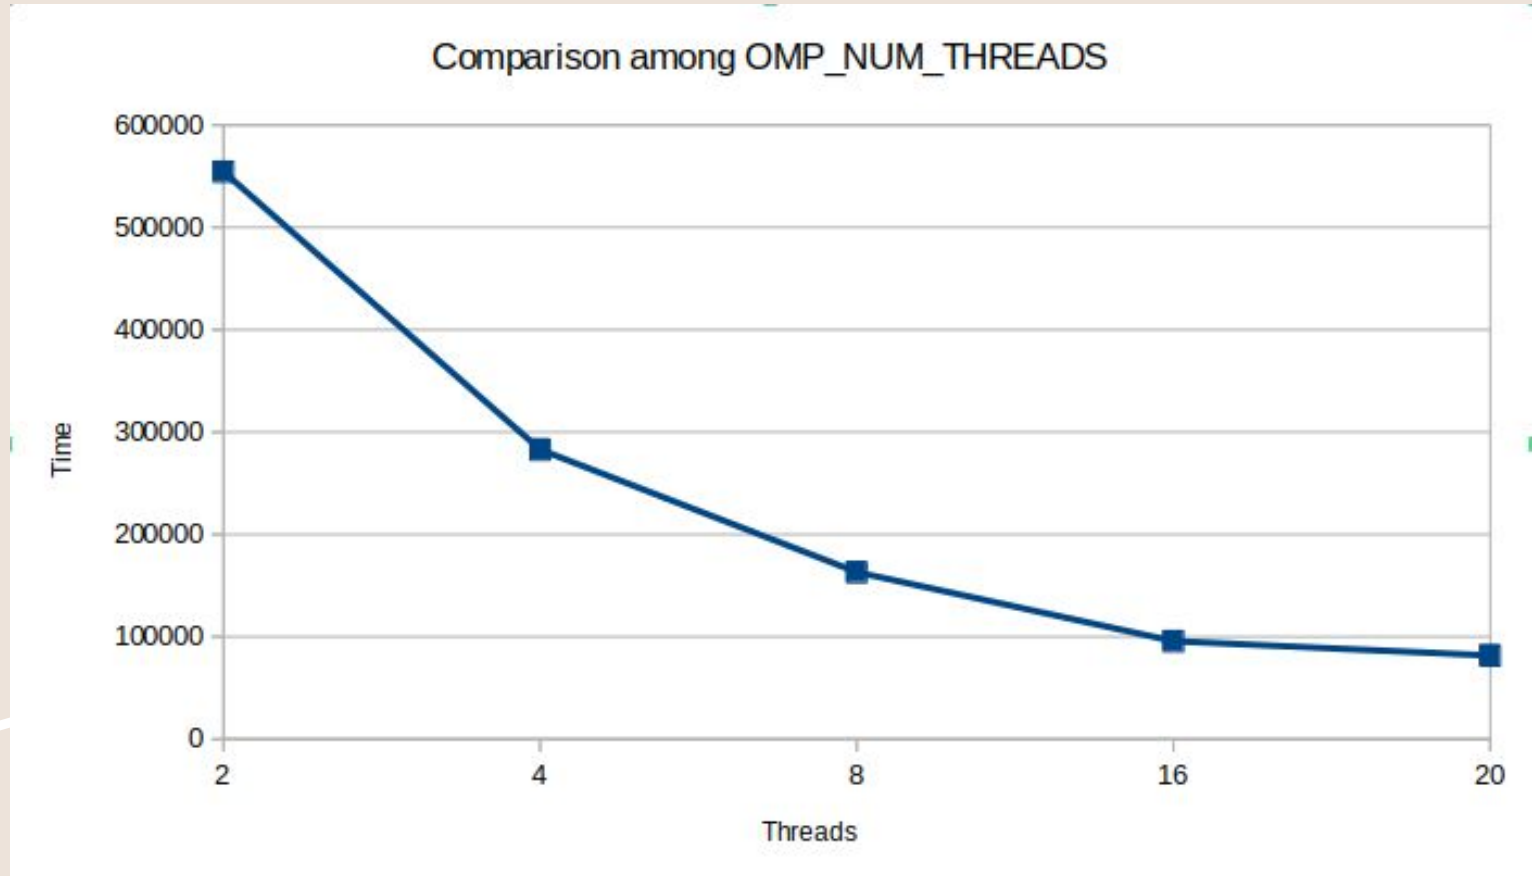


Clearly the Blocked Parallel version takes the lowest time as observed from the graph causing it to be the most optimized algorithm.

# Comparing with fixed matrix size(4096x4096) and different number of nodes for the blocked parallel algorithm

| Number of Threads | Time (ms) |
|-------------------|-----------|
| 2                 | 554371    |
| 4                 | 282853    |
| 8                 | 163085    |
| 16                | 95673     |
| 20                | 81582     |

## Comparing with different number of threads - (fixed matrix size = 4096X4096)



# REFERENCES

- [1] Asghar Aini, Amir Salehipour, "Speeding up the Floyd–Warshall algorithm for the cycled shortest path problem", Applied Mathematics Letters, Volume 25, Issue 1, 2012, Pages 1-5.  
Link: [Speeding up the Floyd–Warshall algorithm for the cycled shortest path problem - ScienceDirect](#)
- [2] Jared Moore and Josh Kalapos, "Floyd-Warshall vs Johnson: Solving All Pairs Shortest Paths in Parallel"  
Link: <https://moorejs.github.io/APSP-in-parallel/>
- [3] Pozder, Nudžejma & Ćorović, Dalila & Herenda, Esma & Divjan, Belmin. (2021). "Towards performance improvement of a parallel Floyd-Warshall algorithm using OpenMP and Intel TBB."  
Link: [Towards performance improvement of a parallel Floyd-Warshall algorithm using OpenMP and Intel TBB](#)